

# Colorado Car Accident Prediction: An Analysis of When and Where an Accident May Occur

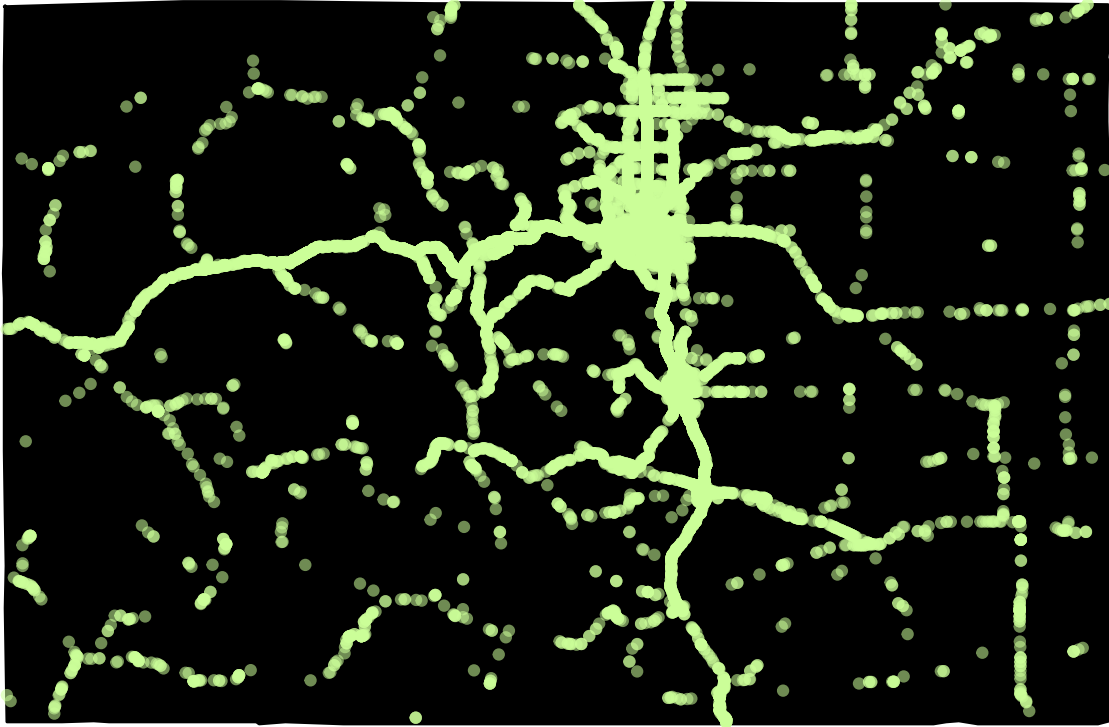
Parker Gauthier

4/17/2022

## Abstract

Automobile accidents are of key concern to many policymakers throughout the country. They should be minimized when possible to prevent substantial economic and social harm. Furthermore, predicting severity can help allocate resources when responding to car accidents. Focused on the state of Colorado, the goal of the following study was to build models predicting the number of car accidents in an area along with the severity of an individual accident. Using data gathered from the Colorado Department of Transportation, this study takes information of over 120,000 auto accidents that took place in the year 2019. This was then used to try and predict, given certain factors, how many accidents will result in injury. The model used in making these predictions was a binomial logistic model using Lasso regularization. With an approximately 79% success rate, this model was marginally better at predicting injury than a model that chose to only predict no injury (approximately 77% of accidents in our test resulted in no injury). Despite this, the coefficients extracted from the model show intuitive factors that can cause injury. In addition to this, another data set was used, gathered from a public data set on Kaggle.com, containing accident counts and different weather conditions for the day. This set was used in conjunction to the other to build a Gradient Boosted Forests model aimed at predicting the number of accidents per 100,000 people for a particular day in specific counties. This model was able to predict with decent accuracy the number of out-of-sample accidents that would happen in a day given particular weather factors.

Figure 1 – Car Accidents in Colorado (2016–2021)



## Introduction

The plot above shows us a sample of 25,340 automobile accidents that have occurred in the state of Colorado. At this level, we can see city centers outlined by highways and interstates where traffic is dense and car accidents are common. Behind each of these accidents are costs ranging from minor property damage to the unfortunate loss of life. The essence of the following analysis will be to determine factors that may have contributed to these accidents and their severity, as well as gauge where and when accidents will occur given certain circumstances. The variables of interest will include weather conditions, location, time of day, and roadway characteristics. By looking at these relationships the hope is to give resources to policy makers and first-responders. If severity can be predicted by characteristics, particularly those that can be reported on a 911 call, EMTs may be able to better gauge the situation they are heading into. Furthermore, if the number of accidents can be predicted given certain weather conditions and times of the year, policymakers may be able to allocate funds appropriately to areas that may need the most relief. If predictions like these are accurate, it could potentially guide policy in a way that improves safety for Colorado drivers.

## Methods and Data

There are three data sets used in conjunction for this study. The first comes from the Colorado Department of Transportation. This data is a record of automobile accidents in the state of Colorado for the year 2019. It includes over 120,000 events with several variables of interest. Dates and times are reported as well as location information providing names of the counties and streets the events took place. Furthermore, there are several variables describing road conditions leading up to an event. These include road type (highway/interstate/city street/urban/rural), weather statuses (Icy/Dry/Wet), proximity to intersections, speed limits (mph), road lighting (Daylight/Dark-unlighted, Dark-lighted), and road contour (Level vs Hilltop).

There are also features describing driver traits such as age (coded 1 for teen, 0 for adult), speed at time of accident (mph), inebriation status (coded 1 for intoxicated, 0 if not), driver behavior prior to the accident (i.e. asleep at the wheel, weaving, changing lanes), and what type of vehicle they were driving (coded 1 for motorcycle, 0 if not). Two other key variable of interest will be the severity of the accident (coded 0 for no injury and 1 for injury/fatality) and the most harmful event (head-on collision, running into a deer, etc).

This data was used to build a binomial logistic regression model using Lasso regularization. This specification was used to predict whether an accident would be classified as non-severe (no injury) or severe (injury or fatality). Lasso regularization was employed to quickly select the features to be used in the regression as there are many in the data used. The outcome variable was regressed on several factors. These included the type of road the accident occurred on (county road vs state highway), the number of vehicles involved in the accident, the condition of the roadway (dry vs icy vs wet), the most harmful event (head-on collision, sideswipe, tree, etc), the speed limit in the area, the speed of the fastest car involved, and driver characteristics (teen, driving motorcycle, impairment status, and if they were wearing a seat belt). This model was trained on 80% of the data using cross validation and then was tested on the remaining 20%. A confusion matrix was then constructed to determine the out-of-sample accuracy. This number was then compared to a model that predicted no-injury for all accidents as a baseline.

Our other data set comes from real-time traffic information collected by researchers Sobhan Moosavi, Mohammad Hossein Samavatian, Srinivasan Parthasarathy, and Rajiv Ramnath. This information was collected by scraping multiple APIs that provide live information on new traffic incidents in the United States. The researchers have made this data publicly available on Kaggle.com to aid the research community in traffic accident analysis. The data set contains nearly three million observations of traffic accidents in the United States between the years 2016-2021. This information includes similar metrics as the previous; however, this data set has latitudinal and longitudinal coordinates as well as more detailed weather information associated with each observation. For this reason, certain visualizations were constructed using this data set and it was helpful in aggregating county-level weather data. Due to its large file size the csv was filtered separately for instances that occurred in Colorado then saved into a new file.

To aggregate county level data, first counts of accidents were taken from the first data set grouped by date and county. This number was then divided by the total population in each county (given by Colorado's Information Marketplace; the third data set) and multiplied by 100,000 to get a population adjusted rate of accidents. Then, using the Kaggle data set, averages of temperatures, wind speeds, precipitation, and visibility were taken for particular dates in different counties. From here, a gradient boosted trees model was constructed. The interaction depth was chosen to be 10 and the optimal number of trees used was selected through cross validation of 10 folds. This model was trained on 80% of the county level data and then was tested on the remaining 20%. Then, the root mean squared error was taken from the models prediction on the out-of-sample outcomes that were in the testing set.

## Results

### Predicting Severity using Lasso Regression

Below displays the results for classifying severity using our regularized binomial logistic model. The ROC curve displayed in Figure 2 shows the relationship between the true positive rate and false positive rate when changing the threshold of how we classify injury based on our regression predictions. We can see that it